

```
# Install common libraries
!pip -q install -U scikit-learn pandas numpy matplotlib tensorflow joblib

# Imports and basic setup
import os, io, zipfile, math
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.linear_model import Ridge
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import joblib
import random

# plotting settings
%matplotlib inline
plt.rcParams['figure.figsize'] = (10,5)

# reproducibility
SEED = 42
np.random.seed(SEED)
random.seed(SEED)
```

```
91.2/91.2 kB 1.2 MB/s eta 0:00:00
62.1/62.1 kB 3.9 MB/s eta 0:00:00
9.5/9.5 MB 38.5 MB/s eta 0:00:00
12.0/12.0 MB 65.7 MB/s eta 0:00:00
16.6/16.6 MB 60.5 MB/s eta 0:00:00
8.7/8.7 MB 70.0 MB/s eta 0:00:00
620.7/620.7 MB 2.9 MB/s eta 0:00:00
5.5/5.5 MB 73.7 MB/s eta 0:00:00
```

ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the following dependency conflicts.
google-colab 1.0.0 requires pandas==2.2.2, but you have pandas 2.3.2 which is incompatible.
opencv-python-headless 4.12.0.88 requires numpy<2.3.0,>=2; python_version >= "3.9", but you have numpy 2.3.2 which is incompatible.
tensorflow-text 2.19.0 requires tensorflow<2.20,>=2.19.0, but you have tensorflow 2.20.0 which is incompatible.
cupy-cuda12x 13.3.0 requires numpy<2.3,>=1.22, but you have numpy 2.3.2 which is incompatible.
opencv-contrib-python 4.12.0.88 requires numpy<2.3.0,>=2; python_version >= "3.9", but you have numpy 2.3.2 which is incompatible.
cudf-cu12 25.6.0 requires pandas<2.2.4dev0,>=2.0, but you have pandas 2.3.2 which is incompatible.
opencv-python 4.12.0.88 requires numpy<2.3.0,>=2; python_version >= "3.9", but you have numpy 2.3.2 which is incompatible.
dask-cudf-cu12 25.6.0 requires pandas<2.2.4dev0,>=2.0, but you have pandas 2.3.2 which is incompatible.
numba 0.60.0 requires numpy<2.1,>=1.22, but you have numpy 2.3.2 which is incompatible.
tf-keras 2.19.0 requires tensorflow<2.20,>=2.19, but you have tensorflow 2.20.0 which is incompatible.
tensorflow-decision-forests 1.12.0 requires tensorflow==2.19.0, but you have tensorflow 2.20.0 which is incompatible.

```
file_path = None # if you already have a path, set it here, e.g. "/content/tsla.csv"
```

```
if file_path is None:
    from google.colab import files
    print("Please upload your CSV or ZIP file (Kaggle dataset).")
    uploaded = files.upload()
    # take the first uploaded filename
    file_path = next(iter(uploaded.keys()))

print("Using file:", file_path)
```

Please upload your CSV or ZIP file (Kaggle dataset).

No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving archive (1).zip to archive (1) (1).zip
Using file: archive (1) (1).zip

```
# Supports .csv or .zip (with a csv inside)
def load_csv_from_path(path):
    path = str(path)
    if path.lower().endswith('.csv'):
        return pd.read_csv(path)
    if path.lower().endswith('.zip'):
        z = zipfile.ZipFile(path)
        # pick first CSV inside
        csvs = [n for n in z.namelist() if n.lower().endswith('.csv')]
        if not csvs:
            raise ValueError("ZIP contains no CSV files.")
        with z.open(csvs[0]) as f:
            return pd.read_csv(f)
        raise ValueError("Please provide a .csv or .zip file.")

df = load_csv_from_path(file_path)
```

```
print("Raw shape:", df.shape)
display(df.head())
print("\nColumns:", list(df.columns))
```

```
Raw shape: (2516, 20)
```

	date	open	high	low	close	volume	rsi_7	rsi_14	cci_7	cci_14	sma_50	ema_50	sma_100
0	2014-01-02	9.986667	10.165333	9.770000	10.006667	92826000	55.344071	54.440118	-37.373644	15.213422	9.682107	9.820167	10.494241
1	2014-01-03	10.000000	10.146000	9.906667	9.970667	70425000	53.742629	53.821521	-81.304471	17.481130	9.652800	9.826069	10.49569:
2	2014-01-06	10.000000	10.026667	9.682667	9.800000	80416500	46.328174	50.870410	-123.427544	-37.824708	9.629467	9.825047	10.496741
3	2014-01-07	9.841333	10.026667	9.683333	9.957333	75511500	53.263037	53.406750	-84.784651	-20.779431	9.597747	9.830235	10.503401
4	2014-01-08	9.923333	10.246667	9.917333	10.085333	92448000	58.368660	55.423026	60.799662	43.570559	9.573240	9.840239	10.511141

```
Columns: ['date', 'open', 'high', 'low', 'close', 'volume', 'rsi_7', 'rsi_14', 'cci_7', 'cci_14', 'sma_50', 'ema_50', 'sma_100', 'en
```

```
# Normalize column names
df.columns = [c.strip().lower() for c in df.columns]

# Find a date column and convert it
date_col = None
for c in ['date', 'datetime', 'time', 'timestamp']:
    if c in df.columns:
        date_col = c
        break

if date_col:
    df[date_col] = pd.to_datetime(df[date_col], errors='coerce')
    df = df.dropna(subset=[date_col]).sort_values(date_col).reset_index(drop=True)
    print("Using date column:", date_col)
else:
    print("No date-like column found. We'll use row order as time order.")

# Required OHLCV columns
required = ['open', 'high', 'low', 'close', 'volume']
missing = [c for c in required if c not in df.columns]
if missing:
    raise ValueError(f"Missing required columns: {missing}. Please provide a dataset with open/high/low/close/volume.")
else:
    print("Found OHLCV columns.")

# Quick summary
print("After cleaning shape:", df.shape)
display(df[required + [date_col] if date_col else required].head())
```

```
Using date column: date
Found OHLCV columns.
After cleaning shape: (2516, 20)
```

	open	high	low	close	volume	date
0	9.986667	10.165333	9.770000	10.006667	92826000	2014-01-02
1	10.000000	10.146000	9.906667	9.970667	70425000	2014-01-03
2	10.000000	10.026667	9.682667	9.800000	80416500	2014-01-06
3	9.841333	10.026667	9.683333	9.957333	75511500	2014-01-07
4	9.923333	10.246667	9.917333	10.085333	92448000	2014-01-08

```
if 'next_day_close' in df.columns:
    df['y'] = df['next_day_close']
    print("Using provided next_day_close column as target 'y'.")
else:
    df['y'] = df['close'].shift(-1)
    print("Created target 'y' as next row's close (close shifted -1).")

# Drop the final row if y is NaN (because we shifted)
df = df.dropna(subset=['y']).reset_index(drop=True)
print("Rows after creating y:", len(df))
```

```

➦ Using provided next_day_close column as target 'y'.
Rows after creating y: 2516

# We'll use lagged OHLCV and any extra indicators that exist in your file.
base_feats = ['open', 'high', 'low', 'close', 'volume']
# find extras present in your file
possible_extras = ['rsi_7', 'rsi_14', 'cci_7', 'cci_14', 'sma_50', 'ema_50', 'sma_100', 'ema_100', 'macd', 'bollinger', 'atr_7', 'atr_14', 'trueran
extras = [c for c in possible_extras if c in df.columns]

print("Extras found (will include as lagged features if present):", extras)

lags = [1,2,3,5] # you can edit this
for lag in lags:
    for c in base_feats + extras:
        df[f'{c}_lag{lag}'] = df[c].shift(lag)

# Drop rows that became NaN due to lagging
df = df.dropna().reset_index(drop=True)
print("Rows after lagging and dropping NaNs:", len(df))

# Build features list (only lag columns)
feature_cols = [c for c in df.columns if any(c.endswith(f'_lag{l}') for l in lags)]
print("Number of features used for classic ML:", len(feature_cols))
print(feature_cols[:30])

➦ Extras found (will include as lagged features if present): ['rsi_7', 'rsi_14', 'cci_7', 'cci_14', 'sma_50', 'ema_50', 'sma_100', 'en
Rows after lagging and dropping NaNs: 2511
Number of features used for classic ML: 72
['open_lag1', 'high_lag1', 'low_lag1', 'close_lag1', 'volume_lag1', 'rsi_7_lag1', 'rsi_14_lag1', 'cci_7_lag1', 'cci_14_lag1', 'sma_5

n = len(df)
train_end = int(n * 0.70)
val_end = int(n * 0.85)

X = df[feature_cols].values
y = df['y'].values

X_train, y_train = X[:train_end], y[:train_end]
X_val, y_val = X[train_end:val_end], y[train_end:val_end]
X_test, y_test = X[val_end:], y[val_end:]

print("Sizes -> train:", len(X_train), "val:", len(X_val), "test:", len(X_test))

➦ Sizes -> train: 1757 val: 377 test: 377

scaler = StandardScaler()
X_train_s = scaler.fit_transform(X_train)
X_val_s = scaler.transform(X_val)
X_test_s = scaler.transform(X_test)

# We compare to a naive rule: predict next-day close = today's close.
close_series = df['close'].values # note: df already had lagged rows removed, so indices align with X/y

# Today's close corresponding to each test row: starting at index val_end in df
y_pred_naive = close_series[val_end: val_end + len(y_test)]

mae_naive = mean_absolute_error(y_test, y_pred_naive)
mse_naive = mean_squared_error(y_test, y_pred_naive)
rmse_naive = math.sqrt(mse_naive)

print("Naive baseline -> MAE: {:.4f}, RMSE: {:.4f}".format(mae_naive, rmse_naive))

➦ Naive baseline -> MAE: 5.7034, RMSE: 7.5539

# Ridge regression (linear)
ridge = Ridge(alpha=1.0)
ridge.fit(X_train_s, y_train)
pred_val_ridge = ridge.predict(X_val_s)

# Random Forest (tree-based)
rf = RandomForestRegressor(n_estimators=300, random_state=SEED, n_jobs=-1)
rf.fit(X_train, y_train) # trees don't require scaling
pred_val_rf = rf.predict(X_val)

# Metric helper (returns MAE, RMSE, R2)

```

```
def metrics(y_true, y_pred):
    mae = mean_absolute_error(y_true, y_pred)
    rmse = math.sqrt(mean_squared_error(y_true, y_pred))
    r2 = r2_score(y_true, y_pred)
    return mae, rmse, r2

print("Ridge val -> MAE, RMSE, R2:", [round(x,4) for x in metrics(y_val, pred_val_ridge)])
print("RF val -> MAE, RMSE, R2:", [round(x,4) for x in metrics(y_val, pred_val_rf)])

# pick best by RMSE on validation
best_model_name = 'rf' if metrics(y_val, pred_val_rf)[1] < metrics(y_val, pred_val_ridge)[1] else 'ridge'
print("Best on validation is:", best_model_name)

➞ Ridge val -> MAE, RMSE, R2: [18.9906, 23.2967, 0.8094]
   RF val -> MAE, RMSE, R2: [53.2133, 71.1206, -0.7768]
   Best on validation is: ridge


# choose model and compute test predictions
if best_model_name == 'rf':
    test_pred = rf.predict(X_test)
else:
    test_pred = ridge.predict(X_test_s)

mae_test, rmse_test, r2_test = metrics(y_test, test_pred)
print("Naive baseline (test) -> MAE: {:.4f}, RMSE: {:.4f}".format(mae_naive, rmse_naive))
print("Best model (test) -> MAE: {:.4f}, RMSE: {:.4f}, R2: {:.4f}".format(mae_test, rmse_test, r2_test))

# Directional accuracy: did the model get up/down right vs previous day's close?
prev_close_for_test = close_series[val_end - 1 : val_end - 1 + len(y_test)] # today's close corresponding to each prediction
true_dir = np.sign(y_test - prev_close_for_test)
pred_dir = np.sign(test_pred - prev_close_for_test)
directional_accuracy = (true_dir == pred_dir).mean()
print("Directional accuracy (up/down): {:.2%}".format(directional_accuracy))

# Plot actual vs predicted (test)
plt.figure()
plt.plot(y_test, label='Actual y (next close)')
plt.plot(test_pred, label='Predicted')
plt.title('Test set - Actual vs Predicted')
plt.legend()
plt.show()

➞ Naive baseline (test) -> MAE: 5.7034, RMSE: 7.5539
   Best model (test) -> MAE: 15.3724, RMSE: 18.8205, R2: 0.8392
   Directional accuracy (up/down): 48.54%
```



```
os.makedirs('models', exist_ok=True)
if best_model_name == 'rf':
    joblib.dump(rf, 'models/price_model_rf.joblib')
else:
    joblib.dump(ridge, 'models/price_model_ridge.joblib')

joblib.dump(scaler, 'models/scaler_standard.joblib')
print("Saved best model and scaler to ./models/")
```

📁 Saved best model and scaler to ./models/

```
import tensorflow as tf
from tensorflow.keras import layers, models, callbacks

# Choose sequence features (raw OHLCV)
seq_feats = ['open', 'high', 'low', 'close', 'volume']
seq_feats = [c for c in seq_feats if c in df.columns] # ensure present

# Build X_all and y_all from the df after lagging (so indexing aligns)
X_all_raw = df[seq_feats].values
y_all = df['y'].values

# Scale features using MinMaxScaler fit on train portion to avoid leakage
train_rows_for_scaler = int(len(X_all_raw) * 0.70)
mm = MinMaxScaler()
mm.fit(X_all_raw[:train_rows_for_scaler])
X_all_s = mm.transform(X_all_raw)

# Create sequences
SEQ_LEN = 60 # last 60 days to predict next day
Xs, ys = [], []
for i in range(len(X_all_s) - SEQ_LEN):
    Xs.append(X_all_s[i : i + SEQ_LEN])
    ys.append(y_all[i + SEQ_LEN])
Xs = np.array(Xs)
ys = np.array(ys)
print("Total sequences:", len(Xs))













































# train/val/test split for sequences (70/15/15 of sequences)
nseq = len(Xs)
t_end = int(nseq * 0.70)
v_end = int(nseq * 0.85)
X_train_nn, y_train_nn = Xs[:t_end], ys[:t_end]
X_val_nn, y_val_nn = Xs[t_end:v_end], ys[t_end:v_end]
X_test_nn, y_test_nn = Xs[v_end:], ys[v_end:]


# Build a small LSTM
tf.random.set_seed(SEED)
model = models.Sequential([
    layers.Input(shape=(SEQ_LEN, X_train_nn.shape[-1])),
    layers.LSTM(64, return_sequences=True),
    layers.Dropout(0.2),
    layers.LSTM(32),
    layers.Dropout(0.2),
    layers.Dense(1)
])
model.compile(optimizer='adam', loss='mse', metrics=['mae'])
es = callbacks.EarlyStopping(patience=8, restore_best_weights=True)


history = model.fit(
    X_train_nn, y_train_nn,
    validation_data=(X_val_nn, y_val_nn),
    epochs=100, batch_size=64,
    callbacks=[es], verbose=1
)


# Evaluate on sequences test set
pred_test_nn = model.predict(X_test_nn).ravel()
mae_l = mean_absolute_error(y_test_nn, pred_test_nn)
rmse_l = math.sqrt(mean_squared_error(y_test_nn, pred_test_nn))
r2_l = r2_score(y_test_nn, pred_test_nn)
print("LSTM test -> MAE: {:.4f}, RMSE: {:.4f}, R2: {:.4f}".format(mae_l, rmse_l, r2_l))


plt.figure()
plt.plot(y_test_nn, label='Actual')
plt.plot(pred_test_nn, label='LSTM Pred')
plt.legend(); plt.title('LSTM - Actual vs Predicted'); plt.show()
```


↩ Total sequences: 2451
Epoch 1/100
27/27  13s 199ms/step - loss: 3126.3525 - mae: 32.2424 - val_loss: 72822.3750 - val_mae: 264.3744
Epoch 2/100
27/27  2s 74ms/step - loss: 2857.4644 - mae: 27.4635 - val_loss: 71634.5078 - val_mae: 262.1181
Epoch 3/100
27/27  3s 80ms/step - loss: 2757.4285 - mae: 25.6469 - val_loss: 71090.2812 - val_mae: 261.0778
Epoch 4/100
27/27  3s 104ms/step - loss: 2709.5369 - mae: 24.6387 - val_loss: 70599.4844 - val_mae: 260.1361
Epoch 5/100
27/27  4s 73ms/step - loss: 2663.3179 - mae: 23.7590 - val_loss: 70104.5859 - val_mae: 259.1831
Epoch 6/100
27/27  2s 70ms/step - loss: 2620.9861 - mae: 22.7725 - val_loss: 69614.2109 - val_mae: 258.2354
Epoch 7/100
27/27  3s 73ms/step - loss: 2580.1692 - mae: 21.8849 - val_loss: 69159.2500 - val_mae: 257.3530
Epoch 8/100
27/27  3s 113ms/step - loss: 2544.6587 - mae: 21.1229 - val_loss: 68725.9062 - val_mae: 256.5096
Epoch 9/100
27/27  4s 70ms/step - loss: 2508.9099 - mae: 20.2540 - val_loss: 68311.3906 - val_mae: 255.7004
Epoch 10/100
27/27  2s 70ms/step - loss: 2469.2856 - mae: 19.7001 - val_loss: 67911.4453 - val_mae: 254.9171
Epoch 11/100
27/27  2s 70ms/step - loss: 2445.3064 - mae: 19.2665 - val_loss: 67527.0234 - val_mae: 254.1620
Epoch 12/100
27/27  2s 71ms/step - loss: 2420.7000 - mae: 18.9608 - val_loss: 67156.2734 - val_mae: 253.4315
Epoch 13/100
27/27  3s 108ms/step - loss: 2393.9177 - mae: 18.6938 - val_loss: 66798.0391 - val_mae: 252.7238
Epoch 14/100
27/27  4s 71ms/step - loss: 2365.5381 - mae: 18.4091 - val_loss: 66452.1250 - val_mae: 252.0385
Epoch 15/100
27/27  2s 71ms/step - loss: 2350.5649 - mae: 18.4740 - val_loss: 66117.3047 - val_mae: 251.3734
Epoch 16/100
27/27  2s 70ms/step - loss: 2326.8669 - mae: 18.3390 - val_loss: 65795.9688 - val_mae: 250.7334
Epoch 17/100
27/27  3s 83ms/step - loss: 2307.8650 - mae: 18.3770 - val_loss: 65482.3359 - val_mae: 250.1072
Epoch 18/100
27/27  3s 98ms/step - loss: 2283.5342 - mae: 18.4240 - val_loss: 65179.2070 - val_mae: 249.5005
Epoch 19/100
27/27  4s 70ms/step - loss: 2279.2205 - mae: 18.6034 - val_loss: 64891.2539 - val_mae: 248.9227
Epoch 20/100
27/27  2s 70ms/step - loss: 2262.6226 - mae: 18.6129 - val_loss: 64611.6289 - val_mae: 248.3604
Epoch 21/100
27/27  2s 71ms/step - loss: 2263.5850 - mae: 18.9320 - val_loss: 64347.2891 - val_mae: 247.8277
Epoch 22/100
27/27  3s 99ms/step - loss: 2240.1714 - mae: 19.0276 - val_loss: 64088.1016 - val_mae: 247.3042
Epoch 23/100
27/27  2s 81ms/step - loss: 2231.2676 - mae: 19.2704 - val_loss: 63839.7734 - val_mae: 246.8017
Epoch 24/100
27/27  2s 70ms/step - loss: 2208.4075 - mae: 19.3696 - val_loss: 63597.8633 - val_mae: 246.3111
Epoch 25/100
27/27  3s 73ms/step - loss: 2209.3381 - mae: 19.6643 - val_loss: 63368.5938 - val_mae: 245.8452
Epoch 26/100
27/27  4s 115ms/step - loss: 2202.0999 - mae: 19.9233 - val_loss: 63148.8242 - val_mae: 245.3978
Epoch 27/100
27/27  3s 112ms/step - loss: 2179.5776 - mae: 20.0809 - val_loss: 62934.5938 - val_mae: 244.9609
Epoch 28/100
27/27  4s 74ms/step - loss: 2168.3672 - mae: 20.2545 - val_loss: 62727.4453 - val_mae: 244.5378
Epoch 29/100
27/27  2s 71ms/step - loss: 2176.3267 - mae: 20.5382 - val_loss: 62525.7891 - val_mae: 244.1251
Epoch 30/100
27/27  3s 74ms/step - loss: 2163.5801 - mae: 20.8442 - val_loss: 62338.5664 - val_mae: 243.7413
Epoch 31/100
27/27  3s 109ms/step - loss: 2173.3140 - mae: 20.9431 - val_loss: 62156.2109 - val_mae: 243.3670
Epoch 32/100
27/27  4s 72ms/step - loss: 2178.2478 - mae: 21.2820 - val_loss: 61989.9453 - val_mae: 243.0252
Epoch 33/100
27/27  3s 76ms/step - loss: 2150.8298 - mae: 21.5132 - val_loss: 61828.2656 - val_mae: 242.6923
Epoch 34/100
27/27  2s 71ms/step - loss: 2143.8394 - mae: 21.6711 - val_loss: 61672.4609 - val_mae: 242.3711
Epoch 35/100
27/27  3s 101ms/step - loss: 2140.9797 - mae: 21.9065 - val_loss: 61523.5508 - val_mae: 242.0637
Epoch 36/100
27/27  2s 83ms/step - loss: 2138.6846 - mae: 22.0612 - val_loss: 61381.0547 - val_mae: 241.7691
Epoch 37/100
27/27  2s 71ms/step - loss: 2137.0989 - mae: 22.4156 - val_loss: 61250.9766 - val_mae: 241.5000
Epoch 38/100
27/27  2s 71ms/step - loss: 2118.4473 - mae: 22.3574 - val_loss: 61121.6562 - val_mae: 241.2321
Epoch 39/100
27/27  3s 72ms/step - loss: 2133.4553 - mae: 22.6293 - val_loss: 61005.1523 - val_mae: 240.9905
Epoch 40/100
27/27  2s 72ms/step - loss: 2134.3049 - mae: 22.9218 - val_loss: 60894.3711 - val_mae: 240.7605
Epoch 41/100
27/27  3s 111ms/step - loss: 2125.7888 - mae: 22.9496 - val_loss: 60789.9883 - val_mae: 240.5437
Epoch 42/100
27/27  2s 72ms/step - loss: 2127.0779 - mae: 23.0469 - val_loss: 60691.3164 - val_mae: 240.3385
Epoch 43/100
27/27  2s 71ms/step - loss: 2122.5417 - mae: 23.1595 - val_loss: 60597.7734 - val_mae: 240.1438
Epoch 44/100
27/27  2s 71ms/step - loss: 2113.7798 - mae: 23.3907 - val_loss: 60508.5156 - val_mae: 239.9579
Epoch 45/100


27/27  2s 70ms/step - loss: 2121.0601 - mae: 23.4877 - val_loss: 60423.9609 - val_mae: 239.7816
Epoch 46/100


27/27  2s 72ms/step - loss: 2126.1050 - mae: 23.6153 - val_loss: 60344.6797 - val_mae: 239.6163
Epoch 47/100


27/27  3s 106ms/step - loss: 2110.5925 - mae: 23.7674 - val_loss: 60269.9180 - val_mae: 239.4602
Epoch 48/100


27/27  4s 71ms/step - loss: 2105.8188 - mae: 23.5810 - val_loss: 60195.3242 - val_mae: 239.3044
Epoch 49/100


27/27  2s 71ms/step - loss: 2116.9258 - mae: 24.0092 - val_loss: 60135.7617 - val_mae: 239.1799
Epoch 50/100


27/27  2s 71ms/step - loss: 2122.7722 - mae: 24.1180 - val_loss: 60080.2617 - val_mae: 239.0639
Epoch 51/100


27/27  2s 81ms/step - loss: 2102.9551 - mae: 24.0811 - val_loss: 60028.4297 - val_mae: 238.9554
Epoch 52/100


27/27  3s 93ms/step - loss: 2121.1887 - mae: 24.1039 - val_loss: 59974.9570 - val_mae: 238.8435
Epoch 53/100


27/27  2s 72ms/step - loss: 2105.0117 - mae: 24.2737 - val_loss: 59926.8047 - val_mae: 238.7427
Epoch 54/100


27/27  2s 72ms/step - loss: 2106.4597 - mae: 24.3091 - val_loss: 59885.5547 - val_mae: 238.6563
Epoch 55/100


27/27  2s 72ms/step - loss: 2116.1565 - mae: 24.3078 - val_loss: 59845.3516 - val_mae: 238.5721
Epoch 56/100


27/27  3s 71ms/step - loss: 2115.3247 - mae: 24.3943 - val_loss: 59809.0586 - val_mae: 238.4960
Epoch 57/100


27/27  4s 116ms/step - loss: 2126.0349 - mae: 24.4777 - val_loss: 59776.7656 - val_mae: 238.4283
Epoch 58/100


27/27  4s 71ms/step - loss: 2105.0586 - mae: 24.2984 - val_loss: 59741.4414 - val_mae: 238.3542
Epoch 59/100


27/27  2s 70ms/step - loss: 2111.3110 - mae: 24.3914 - val_loss: 59711.2891 - val_mae: 238.2910
Epoch 60/100


27/27  3s 71ms/step - loss: 2102.9421 - mae: 24.6123 - val_loss: 59683.4297 - val_mae: 238.2325
Epoch 61/100


27/27  3s 100ms/step - loss: 2101.7192 - mae: 24.4836 - val_loss: 59656.6094 - val_mae: 238.1762
Epoch 62/100


27/27  4s 72ms/step - loss: 2117.0054 - mae: 24.7000 - val_loss: 59634.3242 - val_mae: 238.1294
Epoch 63/100


27/27  2s 72ms/step - loss: 2109.0774 - mae: 24.6398 - val_loss: 59613.5039 - val_mae: 238.0857
Epoch 64/100


27/27  2s 72ms/step - loss: 2107.6777 - mae: 24.5713 - val_loss: 59588.9258 - val_mae: 238.0341
Epoch 65/100


27/27  2s 71ms/step - loss: 2121.8958 - mae: 24.8502 - val_loss: 59574.0508 - val_mae: 238.0028
Epoch 66/100


27/27  3s 103ms/step - loss: 2113.0183 - mae: 24.6644 - val_loss: 59549.6914 - val_mae: 237.9516
Epoch 67/100


27/27  2s 72ms/step - loss: 2110.9082 - mae: 24.8014 - val_loss: 59527.3438 - val_mae: 237.9046
Epoch 68/100


27/27  3s 71ms/step - loss: 2105.9421 - mae: 24.7840 - val_loss: 59516.7070 - val_mae: 237.8823
Epoch 69/100


27/27  2s 73ms/step - loss: 2111.0798 - mae: 24.7521 - val_loss: 59504.1836 - val_mae: 237.8560
Epoch 70/100


27/27  3s 77ms/step - loss: 2109.3071 - mae: 24.8995 - val_loss: 59497.0000 - val_mae: 237.8409
Epoch 71/100


27/27  3s 112ms/step - loss: 2111.3115 - mae: 24.7996 - val_loss: 59481.7656 - val_mae: 237.8089
Epoch 72/100


27/27  4s 72ms/step - loss: 2116.6021 - mae: 24.8825 - val_loss: 59473.7539 - val_mae: 237.7920
Epoch 73/100


27/27  2s 72ms/step - loss: 2122.7771 - mae: 25.1077 - val_loss: 59477.5312 - val_mae: 237.7999
Epoch 74/100


27/27  3s 75ms/step - loss: 2112.3152 - mae: 24.9003 - val_loss: 59471.3789 - val_mae: 237.7870
Epoch 75/100


27/27  3s 99ms/step - loss: 2109.8032 - mae: 24.9065 - val_loss: 59470.3164 - val_mae: 237.7848
Epoch 76/100


27/27  4s 73ms/step - loss: 2112.4058 - mae: 24.8929 - val_loss: 59463.4180 - val_mae: 237.7703
Epoch 77/100


27/27  2s 71ms/step - loss: 2109.9250 - mae: 24.9961 - val_loss: 59461.2383 - val_mae: 237.7657
Epoch 78/100


27/27  3s 74ms/step - loss: 2102.2856 - mae: 24.8825 - val_loss: 59458.6953 - val_mae: 237.7603
Epoch 79/100


27/27  3s 93ms/step - loss: 2097.2256 - mae: 24.9071 - val_loss: 59456.3633 - val_mae: 237.7554
Epoch 80/100


27/27  3s 94ms/step - loss: 2110.1396 - mae: 24.9542 - val_loss: 59453.6016 - val_mae: 237.7496
Epoch 81/100


27/27  5s 75ms/step - loss: 2112.1399 - mae: 24.8563 - val_loss: 59449.9883 - val_mae: 237.7420
Epoch 82/100


27/27  2s 72ms/step - loss: 2103.7795 - mae: 24.9387 - val_loss: 59446.7734 - val_mae: 237.7353
Epoch 83/100


27/27  2s 71ms/step - loss: 2099.2961 - mae: 24.7382 - val_loss: 59593.1641 - val_mae: 238.0433
Epoch 84/100


27/27  3s 96ms/step - loss: 2095.8828 - mae: 24.4156 - val_loss: 59394.3711 - val_mae: 237.6250
Epoch 85/100


27/27  4s 72ms/step - loss: 2117.4319 - mae: 24.7678 - val_loss: 60500.3359 - val_mae: 239.9413
Epoch 86/100

27/27  3s 75ms/step - loss: 2124.8831 - mae: 23.4447 - val_loss: 60391.8750 - val_mae: 239.7220
Epoch 87/100

27/27  3s 81ms/step - loss: 2086.6172 - mae: 23.0157 - val_loss: 59966.9453 - val_mae: 238.8271
Epoch 88/100

27/27  3s 109ms/step - loss: 2085.2852 - mae: 22.4586 - val_loss: 60003.1914 - val_mae: 238.9125
Epoch 89/100

27/27  4s 76ms/step - loss: 1970.9771 - mae: 17.0206 - val_loss: 59237.8906 - val_mae: 237.2963
Epoch 90/100

27/27  2s 72ms/step - loss: 1914.1005 - mae: 15.6450 - val_loss: 58287.2250 - val_mae: 225.2847
Epoch 91/100

```
Epoch 91/100 2s 72ms/step - loss: 1875.8096 - mae: 15.0428 - val_loss: 57632.3203 - val_mae: 233.8883
27/27
Epoch 92/100 3s 103ms/step - loss: 1806.9805 - mae: 14.7875 - val_loss: 57065.9453 - val_mae: 232.6743
27/27
Epoch 93/100 4s 72ms/step - loss: 1800.7455 - mae: 14.5696 - val_loss: 56523.5273 - val_mae: 231.5057
27/27
Epoch 94/100 3s 71ms/step - loss: 1768.5190 - mae: 14.3590 - val_loss: 56002.0859 - val_mae: 230.3768
27/27
Epoch 95/100 2s 72ms/step - loss: 1767.4841 - mae: 14.1795 - val_loss: 55502.3047 - val_mae: 229.2895
27/27
Epoch 96/100 3s 96ms/step - loss: 1718.4257 - mae: 14.0758 - val_loss: 55012.1562 - val_mae: 228.2181
27/27
Epoch 97/100 2s 88ms/step - loss: 1700.0615 - mae: 14.0128 - val_loss: 54540.1953 - val_mae: 227.1817
27/27
Epoch 98/100 2s 71ms/step - loss: 1671.2545 - mae: 13.7754 - val_loss: 54081.0547 - val_mae: 226.1690
27/27
Epoch 99/100 3s 72ms/step - loss: 1657.9788 - mae: 13.5133 - val_loss: 53635.5508 - val_mae: 225.1819
27/27
Epoch 100/100 2s 72ms/step - loss: 1627.7742 - mae: 13.3201 - val_loss: 53199.0000 - val_mae: 224.2105
12/12 1s 41ms/step
LSTM test -> MAE: 177.7687, RMSE: 183.9896, R2: -14.0417
```

LSTM — Actual vs Predicted

